

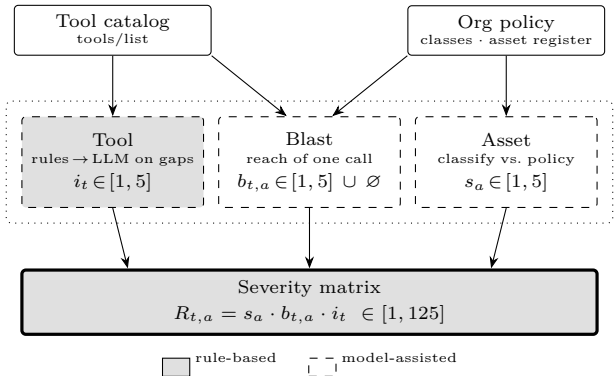


Fig. 1. Static (design-time) scoring. Two documents—the server’s advertised tool catalog and the organization’s written policy—feed three scanners, each producing exactly one primitive: tool impact i_t over the T tools, asset sensitivity s_a over the A assets, and blast radius $b_{t,a}$ over all $T \times A$ pairs. Their product is the severity matrix. Since i and s are each indexed by a single dimension, b is what distinguishes a cell from others in its row; $b_{t,a} = \emptyset$ marks a tool that cannot reach an asset and removes the cell. The organization states no severity number: s_a is derived by classifying each asset against the policy’s own classes.